

Data Comp - Team 322

April 5, 2020

1 Introduction

The New York Police Department, or NYPD, is one of the most active police departments in the US, where they go through many conflicts with criminals throughout the state with a multitude of different crimes ranging from murder, marijuana possession, robbery, arson, etc. Through many instances of arrests in the state, we can see some specific instances of discrimination of specific arrests where the perpetrator can be wrongly accused or wrongly acquitted from their crimes.

This project addresses the following question: Is there evidence of discrimination in NYPD arrests based on demographics?

The approach used will use machine learning algorithms to find insights into how the different variables in the given dataset of NYPD arrests interact with each other. In particular, our approach will use Random Forests to classify arrests based on various predictors.

2 Data and Problem Definition

2.1 Data Set

The given dataset contains variables of each arrest such as the age group, gender, precinct (city district), race, and other descriptive factors. In particular, each row contains information about the arrest key, arrest date, internal classification code, description of said code, a more general internal classification code, description of said code, law code charges, level of offense committed, the borough the arrest was made in, precinct where the arrest occurred, jurisdiction responsible for arrest, age group of the perpetrator, sex of the perpetrator, race of the perpetrator, and coordinates of the location of arrest.

In this project, we mainly use the following features: internal classification code, level of offense committed, age group of the perpetrator, sex of the perpetrator, and race of the perpetrator.

2.2 Problem Definition

In this particular approach, the problem is to find bias in classification in favor of or against a certain demographic.

Our approach will be as follows: Train a model using the data to predict a certain characteristic of the arrest and compare the predictions with the actual labels to see if the predictions are skewed.

3 Methodology


3.1 Exploratory Data Analysis


```
In [20]: import pandas as pd
        nypd = pd.read_csv("NYPD_Arrests_Data__Historic_.csv")


In [21]: nypd = nypd[['PD_DESC', 'LAW_CAT_CD', 'AGE_GROUP', 'PERP_SEX', 'PERP_RACE', 'ARREST_B
        nypd.dropna(inplace=True)

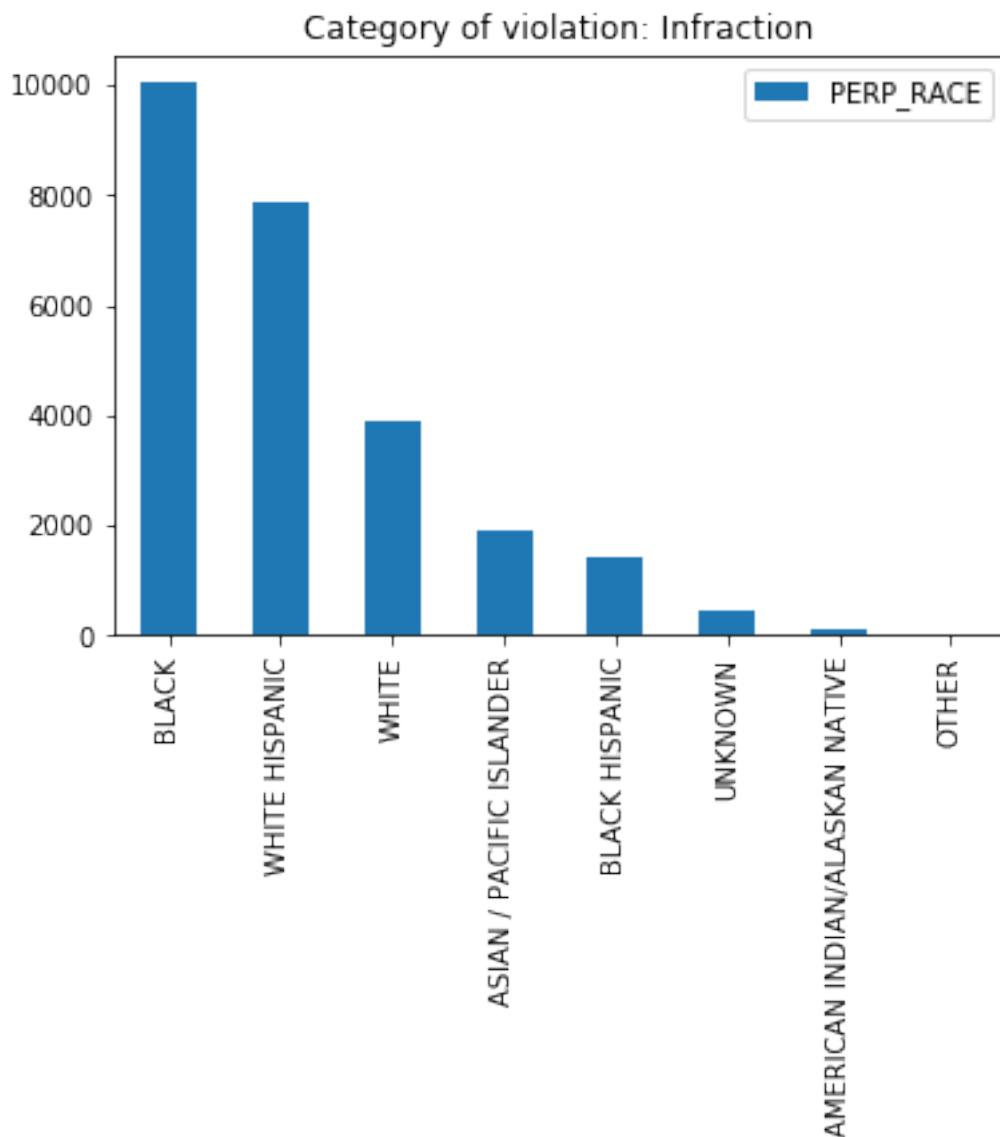
In [22]: dict = {'V': 'Violation', 'F': 'Felony', 'M': 'Misdemeanor', 'I': 'Infraction'}
        for i in nypd.LAW_CAT_CD.unique():
            group = nypd[nypd.LAW_CAT_CD == i]
            counts = group.PERP_RACE.value_counts()
            counts = pd.DataFrame(counts)
            title = "Category of violation: " + dict[i]
            print(counts.plot(kind='bar', title = title))

AxesSubplot(0.125,0.125;0.775x0.755)
AxesSubplot(0.125,0.125;0.775x0.755)
AxesSubplot(0.125,0.125;0.775x0.755)
AxesSubplot(0.125,0.125;0.775x0.755)
```









Using clustering, we get the following output:

Standardized Cluster Means: centroid pd_desc perp_race age_group arrest_boro perp_sex law_cat_cd -----

```
-- 1 ASSAULT 3 BLACK 25-44 Q M M 2 THEFT OF SERVICES, UNCLASSIFIED BLACK
18-24 M M M 3 CONTROLLED SUBSTANCE, POSSESSION 7 WHITE HISPANIC 25-44 K
M M 4 CONTROLLED SUBSTANCE,SALE 3 WHITE HISPANIC 25-44 M M F 5 NY STATE
LAWS,UNCLASSIFIED VIOLATION BLACK 25-44 M M V 6 ASSAULT 2,1,UNCLASSIFIED
WHITE HISPANIC 25-44 B F F
```

Centroid Statistics: centroid size within_cluster_sum_of_squares -- -----

```
----- 1 2.03914e+06 5.26579e+06 2 1.17129e+06 2.77919e+06 3 722192 1.5601e+06 4 416880
969848 5 235334 393414 6 213506 546798
```

3.2 Modeling

We start by using a Naive Bayes Classifier to predict the race of the perpetrator using their age group, sex, category of offense, and description of offense. Even though this is counter intuitive, we are trying to uncover a pattern that the model could learn about the race of the perpetrator in specific arrest scenarios.

```
In [32]: from sklearn.model_selection import train_test_split
import h2o
from h2o.estimators.naive_bayes import H2ONaiveBayesEstimator

h2o.init(max_mem_size = "10G")           #specify max number of bytes. uses all cores
h2o.remove_all()

data=nypd[['PD_DESC', 'AGE_GROUP', 'PERP_SEX', 'LAW_CAT_CD', 'PERP_RACE']]

predictors = data.columns.values[:-1]
target = data.columns.values[-1]

train, test = train_test_split(nypd, test_size=0.3)

h_train = h2o.H2OFrame(train)
h_test = h2o.H2OFrame(test)

model = H2ONaiveBayesEstimator(score_each_iteration=True,
                                seed = 1234)
model.train(training_frame = h_train,
            validation_frame = h_test,
            y = str(target),
            x = list(predictors))
print(model)
```

Checking whether there is an H2O instance running at <http://localhost:54321> . connected.

```
-----
H2O_cluster_uptime:      25 mins 51 secs
H2O_cluster_timezone:    America/Los_Angeles
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.30.0.1
H2O_cluster_version_age: 1 day
H2O_cluster_name:        H2O_from_python_varch_j63zsy
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 9.97 Gb
H2O_cluster_total_cores: 4
H2O_cluster_allowed_cores: 4
H2O_cluster_status:      locked, healthy
```

```

H2O_connection_url:      http://localhost:54321
H2O_connection_proxy:    {"http": null, "https": null}
H2O_internal_security:   False
H2O_API_Extensions:      Amazon S3, Algos, AutoML, Core V3, TargetEncoder, Core V4
Python_version:          3.7.1 final
-----

```

```

Parse progress: || 100%
Parse progress: || 100%
naivebayes Model Build progress: || 100%

```

Model Details

=====

H2ONaiveBayesEstimator : Naive Bayes

Model Key: NaiveBayes_model_python_1586062955647_1

Model Summary:

```

      number_of_response_levels  min_apriori_probability  \
0                               8.0                    0.000284

```

```

      max_apriori_probability
0                          0.485656

```

ModelMetricsMultinomial: naivebayes

** Reported on train data. **

MSE: 0.4683722936647309

RMSE: 0.684377303586794

LogLoss: 1.3171315620859425

Mean Per-Class Error: 0.8687579092919805

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

```

      AMERICAN INDIAN/ALASKAN NATIVE  ASIAN / PACIFIC ISLANDER  BLACK  \
0                                0.0                    14.0    6884.0
1                                2.0                    724.0   122904.0
2                                2.0                    502.0  1588126.0
3                                0.0                    200.0   256613.0
4                                0.0                     1.0    907.0
5                                1.0                     60.0   31237.0
6                                3.0                    364.0  373104.0
7                                1.0                    663.0  822796.0

```


8 9.0 2528.0 3202571.0

| | BLACK | HISPANIC | OTHER | UNKNOWN | WHITE | WHITE HISPANIC | Error \ |
|---|-------|----------|-------|---------|---------|----------------|----------|
| 0 | | 2.0 | 0.0 | 2.0 | 75.0 | 380.0 | 1.000000 |
| 1 | | 30.0 | 0.0 | 63.0 | 1776.0 | 9264.0 | 0.994628 |
| 2 | | 110.0 | 0.0 | 283.0 | 6639.0 | 27147.0 | 0.021372 |
| 3 | | 0.0 | 0.0 | 73.0 | 1378.0 | 9221.0 | 1.000000 |
| 4 | | 0.0 | 0.0 | 0.0 | 5.0 | 35.0 | 1.000000 |
| 5 | | 7.0 | 0.0 | 0.0 | 485.0 | 2301.0 | 1.000000 |
| 6 | | 108.0 | 0.0 | 165.0 | 7767.0 | 21939.0 | 0.980749 |
| 7 | | 75.0 | 0.0 | 236.0 | 6164.0 | 40643.0 | 0.953315 |
| 8 | | 332.0 | 0.0 | 822.0 | 24289.0 | 110930.0 | 0.510020 |

| | Rate |
|---|-----------------------|
| 0 | 7,357 / 7,357 |
| 1 | 134,039 / 134,763 |
| 2 | 34,683 / 1,622,809 |
| 3 | 267,485 / 267,485 |
| 4 | 948 / 948 |
| 5 | 34,091 / 34,091 |
| 6 | 395,683 / 403,450 |
| 7 | 829,935 / 870,578 |
| 8 | 1,704,221 / 3,341,481 |

Top-8 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.489980 |
| 1 | 2 | 0.746147 |
| 2 | 3 | 0.877586 |
| 3 | 4 | 0.951860 |
| 4 | 5 | 0.987473 |
| 5 | 6 | 0.997393 |
| 6 | 7 | 0.999696 |
| 7 | 8 | 1.000000 |

ModelMetricsMultinomial: naivebayes

** Reported on validation data. **

MSE: 0.46793905490871296

RMSE: 0.6840607099583436

LogLoss: 1.3156647265511612

Mean Per-Class Error: 0.8687512321177057

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | AMERICAN INDIAN/ALASKAN NATIVE | ASIAN / PACIFIC ISLANDER | BLACK \ |
|---|--------------------------------|--------------------------|-----------|
| 0 | 0.0 | 4.0 | 2853.0 |
| 1 | 0.0 | 313.0 | 52359.0 |
| 2 | 3.0 | 196.0 | 681506.0 |
| 3 | 0.0 | 71.0 | 109710.0 |
| 4 | 0.0 | 1.0 | 389.0 |
| 5 | 0.0 | 24.0 | 13433.0 |
| 6 | 1.0 | 153.0 | 160004.0 |
| 7 | 0.0 | 253.0 | 352407.0 |
| 8 | 4.0 | 1015.0 | 1372661.0 |

| | BLACK HISPANIC | OTHER | UNKNOWN | WHITE | WHITE HISPANIC | Error \ |
|---|----------------|-------|---------|---------|----------------|----------|
| 0 | 0.0 | 0.0 | 1.0 | 36.0 | 178.0 | 1.000000 |
| 1 | 7.0 | 0.0 | 25.0 | 785.0 | 3960.0 | 0.994552 |
| 2 | 39.0 | 0.0 | 133.0 | 2900.0 | 11692.0 | 0.021484 |
| 3 | 7.0 | 0.0 | 32.0 | 575.0 | 4021.0 | 0.999939 |
| 4 | 0.0 | 0.0 | 0.0 | 1.0 | 24.0 | 1.000000 |
| 5 | 2.0 | 0.0 | 5.0 | 197.0 | 973.0 | 0.999658 |
| 6 | 40.0 | 0.0 | 83.0 | 3218.0 | 9171.0 | 0.981363 |
| 7 | 30.0 | 0.0 | 90.0 | 2635.0 | 17523.0 | 0.953014 |
| 8 | 125.0 | 0.0 | 369.0 | 10347.0 | 47542.0 | 0.509399 |

| | Rate |
|---|---------------------|
| 0 | 3,072 / 3,072 |
| 1 | 57,136 / 57,449 |
| 2 | 14,963 / 696,469 |
| 3 | 114,409 / 114,416 |
| 4 | 415 / 415 |
| 5 | 14,629 / 14,634 |
| 6 | 169,452 / 172,670 |
| 7 | 355,415 / 372,938 |
| 8 | 729,491 / 1,432,063 |

Top-8 Hit Ratios:

| k | hit_ratio |
|-----|-----------|
| 0 1 | 0.490601 |
| 1 2 | 0.746526 |
| 2 3 | 0.878012 |
| 3 4 | 0.952205 |
| 4 5 | 0.987490 |

```

5 6    0.997445
6 7    0.999695
7 8    1.000000

```

Judging the model performance, we conclude that there is no clear pattern in the data with respect to race.

Next, we train a Naive Bayes Classifier on all demographic categorical variables to predict the level of offense.

```

In [35]: h2o.init(max_mem_size = "10G")           #specify max number of bytes. uses all cor
         h2o.remove_all()

data=nypd[['AGE_GROUP', 'PERP_SEX', 'PERP_RACE', 'LAW_CAT_CD']]

predictors = data.columns.values[:-1]
target = data.columns.values[-1]

train, test = train_test_split(nypd, test_size=0.3)

h_train = h2o.H2OFrame(train)
h_test = h2o.H2OFrame(test)

model = H2ONaiveBayesEstimator(score_each_iteration=True,
                               seed = 1234)
model.train(training_frame = h_train,
            validation_frame = h_test,
            y = str(target),
            x = list(predictors))
print(model)

```

Checking whether there is an H2O instance running at <http://localhost:54321> . connected.

```

-----
H2O_cluster_uptime:      51 mins 56 secs
H2O_cluster_timezone:    America/Los_Angeles
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.30.0.1
H2O_cluster_version_age: 1 day
H2O_cluster_name:        H2O_from_python_varch_j63zsy
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 9.97 Gb
H2O_cluster_total_cores: 4

```

```

H2O_cluster_allowed_cores: 4
H2O_cluster_status: locked, healthy
H2O_connection_url: http://localhost:54321
H2O_connection_proxy: {"http": null, "https": null}
H2O_internal_security: False
H2O_API_Extensions: Amazon S3, Algos, AutoML, Core V3, TargetEncoder, Core V4
Python_version: 3.7.1 final
-----

```

```

Parse progress: || 100%
Parse progress: || 100%
naivebayes Model Build progress: || 100%
Model Details
=====

```

```

H2ONaiveBayesEstimator : Naive Bayes
Model Key: NaiveBayes_model_python_1586062955647_3

```

Model Summary:

```

      number_of_response_levels  min_apriori_probability  \
0                               4.0                    0.005317

      max_apriori_probability
0                    0.671163

```

```

ModelMetricsMultinomial: naivebayes
** Reported on train data. **

```

```

MSE: 0.27267301688953166
RMSE: 0.5221810192735195
LogLoss: 0.8098425328356452
Mean Per-Class Error: 0.750008137592646

```

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|------|-----|-----------|------|----------|-----------------------|
| 0 | 0.0 | 5.0 | 878035.0 | 15.0 | 1.000000 | 878,055 / 878,055 |
| 1 | 1.0 | 0.0 | 17767.0 | 0.0 | 1.000000 | 17,768 / 17,768 |
| 2 | 45.0 | 3.0 | 2242605.0 | 25.0 | 0.000033 | 73 / 2,242,678 |
| 3 | 2.0 | 1.0 | 202974.0 | 0.0 | 1.000000 | 202,977 / 202,977 |
| 4 | 48.0 | 9.0 | 3341381.0 | 40.0 | 0.328858 | 1,098,873 / 3,341,478 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.671142 |
| 1 | 2 | 0.933906 |
| 2 | 3 | 0.994660 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: naivebayes

** Reported on validation data. **

MSE: 0.2724527333255598

RMSE: 0.5219700502189372

LogLoss: 0.8097145524908334

Mean Per-Class Error: 0.7500033540457705

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|------|-----|-----------|------|----------|---------------------|
| 0 | 0.0 | 0.0 | 375656.0 | 2.0 | 1.000000 | 375,658 / 375,658 |
| 1 | 0.0 | 0.0 | 7875.0 | 0.0 | 1.000000 | 7,875 / 7,875 |
| 2 | 13.0 | 3.0 | 961820.0 | 8.0 | 0.000025 | 24 / 961,844 |
| 3 | 0.0 | 1.0 | 86684.0 | 1.0 | 0.999988 | 86,685 / 86,686 |
| 4 | 13.0 | 4.0 | 1432035.0 | 11.0 | 0.328367 | 470,242 / 1,432,063 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.671633 |
| 1 | 2 | 0.933949 |
| 2 | 3 | 0.994491 |
| 3 | 4 | 1.000000 |

From the results, we notice that the model classifies almost all instances as misdemeanors, indicating that it is simply favoring the most frequent level rather than exposing any structure within the data. This leads us to the conclusion that demographic variables do not accurately model the level of offense.

3.2.1 Random Forest

We now use a more sophisticated non-parametric model to use a selection of variables to predict the the level of offense. In this, we use a methodology of hiding the borough variable from the model after subsetting the data to a specific borough. Effectively, this results in implicitly borough-specific model. We then proceed to test each model against data from different boroughs to check if there is a bias in the predictions of the levels of offense in any model against any different borough

```
In [37]: from h2o.estimators.gbm import H2OGradientBoostingEstimator
         from h2o.estimators.random_forest import H2ORandomForestEstimator

         boroughs = nypd['ARREST_BORO'].unique()

         nypd = nypd[['PD_DESC', 'PERP_RACE', 'AGE_GROUP', 'ARREST_BORO', 'PERP_SEX', 'LAW_CAT',

         for borough in boroughs:
             df_borough = nypd.loc[nypd['ARREST_BORO'] == borough]
             df_borough.drop(['ARREST_BORO'], axis=1, inplace=True)

             train, test = train_test_split(df_borough, test_size=0.2)

             predictors = train.columns.values[:-1]
             target = train.columns.values[-1]

             model = H2ORandomForestEstimator(
                 model_id="rf_"+borough,
                 ntrees=200,
                 stopping_rounds=7,
                 score_each_iteration=True,
                 max_depth=50,
                 seed=1000000)

             h_train = h2o.H2OFrame(train)
             h_test = h2o.H2OFrame(test)

             model.train(list(predictors), str(target), training_frame=h_train, validation_frame=h_test)
             print(model)

             h2o.save_model(model=model, path="rf_"+borough, force=True)
             print(model.score_history())

Parse progress: || 100%
Parse progress: || 100%
drf Model Build progress: || 100%
Model Details
=====
H2ORandomForestEstimator : Distributed Random Forest
Model Key: rf_Q
```

Model Summary:

| | number_of_trees | number_of_internal_trees | model_size_in_bytes | \ |
|---|-----------------|--------------------------|---------------------|---|
| 0 | 27.0 | 108.0 | 386750.0 | |

| | min_depth | max_depth | mean_depth | min_leaves | max_leaves | mean_leaves |
|---|-----------|-----------|------------|------------|------------|-------------|
| 0 | 8.0 | 16.0 | 12.148149 | 28.0 | 298.0 | 133.24074 |

ModelMetricsMultinomial: drf

** Reported on train data. **

MSE: 0.013171771670983146

RMSE: 0.11476833914883994

LogLoss: 0.05353537307174825

Mean Per-Class Error: 0.14278233045171151

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|----------|--------|----------|---------|----------|------------------|
| 0 | 192838.0 | 0.0 | 4514.0 | 0.0 | 0.022873 | 4,514 / 197,352 |
| 1 | 0.0 | 4411.0 | 5176.0 | 143.0 | 0.546660 | 5,319 / 9,730 |
| 2 | 531.0 | 17.0 | 457625.0 | 0.0 | 0.001196 | 548 / 458,173 |
| 3 | 0.0 | 0.0 | 19.0 | 47408.0 | 0.000401 | 19 / 47,427 |
| 4 | 193369.0 | 4428.0 | 467334.0 | 47551.0 | 0.014593 | 10,400 / 712,682 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.985407 |
| 1 | 2 | 0.999572 |
| 2 | 3 | 0.999990 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: drf

** Reported on validation data. **

MSE: 0.013276702073372752

RMSE: 0.11522457235057439

LogLoss: 0.0543130314762207

Mean Per-Class Error: 0.13920517395963944

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|---------|--------|----------|---------|----------|-----------------|
| 0 | 48443.0 | 0.0 | 1199.0 | 0.0 | 0.024153 | 1,199 / 49,642 |
| 1 | 0.0 | 1174.0 | 1287.0 | 42.0 | 0.530963 | 1,329 / 2,503 |
| 2 | 130.0 | 7.0 | 113998.0 | 0.0 | 0.001200 | 137 / 114,135 |
| 3 | 0.0 | 0.0 | 6.0 | 11885.0 | 0.000505 | 6 / 11,891 |
| 4 | 48573.0 | 1181.0 | 116490.0 | 11927.0 | 0.014991 | 2,671 / 178,171 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.985009 |
| 1 | 2 | 0.999506 |
| 2 | 3 | 0.999977 |
| 3 | 4 | 1.000000 |

Scoring History:

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|------------|-----------------|-----------------|
| 0 | 2020-04-04 23:04:49 | 0.112 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:04:53 | 3.980 sec | 1.0 | 0.163007 |
| 2 | 2020-04-04 23:04:56 | 6.781 sec | 2.0 | 0.138292 |
| 3 | 2020-04-04 23:04:58 | 8.925 sec | 3.0 | 0.129223 |
| 4 | 2020-04-04 23:05:00 | 10.854 sec | 4.0 | 0.124490 |
| 5 | 2020-04-04 23:05:02 | 12.713 sec | 5.0 | 0.121484 |
| 6 | 2020-04-04 23:05:04 | 15.035 sec | 6.0 | 0.119587 |
| 7 | 2020-04-04 23:05:06 | 17.313 sec | 7.0 | 0.118227 |
| 8 | 2020-04-04 23:05:09 | 19.574 sec | 8.0 | 0.117259 |
| 9 | 2020-04-04 23:05:11 | 21.875 sec | 9.0 | 0.116538 |
| 10 | 2020-04-04 23:05:13 | 24.394 sec | 10.0 | 0.115940 |
| 11 | 2020-04-04 23:05:16 | 26.951 sec | 11.0 | 0.115600 |
| 12 | 2020-04-04 23:05:19 | 29.730 sec | 12.0 | 0.115542 |
| 13 | 2020-04-04 23:05:22 | 32.544 sec | 13.0 | 0.115182 |
| 14 | 2020-04-04 23:05:24 | 35.355 sec | 14.0 | 0.114962 |
| 15 | 2020-04-04 23:05:28 | 38.541 sec | 15.0 | 0.115219 |
| 16 | 2020-04-04 23:05:31 | 41.776 sec | 16.0 | 0.114949 |
| 17 | 2020-04-04 23:05:34 | 45.065 sec | 17.0 | 0.114733 |
| 18 | 2020-04-04 23:05:38 | 48.910 sec | 18.0 | 0.114543 |

19 2020-04-04 23:05:42 52.722 sec 19.0 0.114392

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.124956 | 0.015037 | 0.163222 |
| 2 | 0.088431 | 0.014741 | 0.132450 |
| 3 | 0.075707 | 0.014691 | 0.123747 |
| 4 | 0.068961 | 0.014758 | 0.120153 |
| 5 | 0.063396 | 0.014701 | 0.118261 |
| 6 | 0.060611 | 0.014688 | 0.117189 |
| 7 | 0.058778 | 0.014663 | 0.116513 |
| 8 | 0.057572 | 0.014613 | 0.116110 |
| 9 | 0.055820 | 0.014637 | 0.115756 |
| 10 | 0.054638 | 0.014632 | 0.115487 |
| 11 | 0.054029 | 0.014639 | 0.115343 |
| 12 | 0.053913 | 0.014633 | 0.115329 |
| 13 | 0.053047 | 0.014637 | 0.115166 |
| 14 | 0.052763 | 0.014629 | 0.115084 |
| 15 | 0.053969 | 0.014625 | 0.115265 |
| 16 | 0.053403 | 0.014623 | 0.115148 |
| 17 | 0.052874 | 0.014616 | 0.115046 |
| 18 | 0.052404 | 0.014608 | 0.114957 |
| 19 | 0.052007 | 0.014606 | 0.114887 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.125809 | 0.015210 |
| 2 | 0.083278 | 0.015008 |
| 3 | 0.070831 | 0.014991 |
| 4 | 0.064559 | 0.014980 |
| 5 | 0.060747 | 0.014986 |
| 6 | 0.058372 | 0.014963 |
| 7 | 0.056646 | 0.014980 |
| 8 | 0.056152 | 0.014963 |
| 9 | 0.055038 | 0.014986 |
| 10 | 0.054155 | 0.014986 |
| 11 | 0.053803 | 0.014991 |
| 12 | 0.053995 | 0.014991 |
| 13 | 0.053385 | 0.014991 |
| 14 | 0.053170 | 0.014986 |
| 15 | 0.054281 | 0.014986 |
| 16 | 0.053809 | 0.014986 |
| 17 | 0.053352 | 0.014991 |
| 18 | 0.052935 | 0.014986 |
| 19 | 0.052648 | 0.014986 |

See the whole table with `table.as_data_frame()`

Variable Importances:

| | variable | relative_importance | scaled_importance | percentage |
|---|-----------|---------------------|-------------------|------------|
| 0 | PD_DESC | 6.984099e+06 | 1.000000 | 0.995673 |
| 1 | PERP_RACE | 1.877215e+04 | 0.002688 | 0.002676 |
| 2 | PERP_SEX | 7.393794e+03 | 0.001059 | 0.001054 |
| 3 | AGE_GROUP | 4.187105e+03 | 0.000600 | 0.000597 |

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|------------------|-----------------|-----------------|
| 0 | 2020-04-04 23:04:49 | 0.112 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:04:53 | 3.980 sec | 1.0 | 0.163007 |
| 2 | 2020-04-04 23:04:56 | 6.781 sec | 2.0 | 0.138292 |
| 3 | 2020-04-04 23:04:58 | 8.925 sec | 3.0 | 0.129223 |
| 4 | 2020-04-04 23:05:00 | 10.854 sec | 4.0 | 0.124490 |
| 5 | 2020-04-04 23:05:02 | 12.713 sec | 5.0 | 0.121484 |
| 6 | 2020-04-04 23:05:04 | 15.035 sec | 6.0 | 0.119587 |
| 7 | 2020-04-04 23:05:06 | 17.313 sec | 7.0 | 0.118227 |
| 8 | 2020-04-04 23:05:09 | 19.574 sec | 8.0 | 0.117259 |
| 9 | 2020-04-04 23:05:11 | 21.875 sec | 9.0 | 0.116538 |
| 10 | 2020-04-04 23:05:13 | 24.394 sec | 10.0 | 0.115940 |
| 11 | 2020-04-04 23:05:16 | 26.951 sec | 11.0 | 0.115600 |
| 12 | 2020-04-04 23:05:19 | 29.730 sec | 12.0 | 0.115542 |
| 13 | 2020-04-04 23:05:22 | 32.544 sec | 13.0 | 0.115182 |
| 14 | 2020-04-04 23:05:24 | 35.355 sec | 14.0 | 0.114962 |
| 15 | 2020-04-04 23:05:28 | 38.541 sec | 15.0 | 0.115219 |
| 16 | 2020-04-04 23:05:31 | 41.776 sec | 16.0 | 0.114949 |
| 17 | 2020-04-04 23:05:34 | 45.065 sec | 17.0 | 0.114733 |
| 18 | 2020-04-04 23:05:38 | 48.910 sec | 18.0 | 0.114543 |
| 19 | 2020-04-04 23:05:42 | 52.722 sec | 19.0 | 0.114392 |
| 20 | 2020-04-04 23:05:46 | 56.899 sec | 20.0 | 0.114302 |
| 21 | 2020-04-04 23:05:50 | 1 min 1.436 sec | 21.0 | 0.114465 |
| 22 | 2020-04-04 23:05:55 | 1 min 6.024 sec | 22.0 | 0.115353 |
| 23 | 2020-04-04 23:06:00 | 1 min 10.781 sec | 23.0 | 0.115177 |
| 24 | 2020-04-04 23:06:05 | 1 min 15.644 sec | 24.0 | 0.115007 |
| 25 | 2020-04-04 23:06:10 | 1 min 20.825 sec | 25.0 | 0.114877 |
| 26 | 2020-04-04 23:06:15 | 1 min 25.855 sec | 26.0 | 0.114902 |
| 27 | 2020-04-04 23:06:20 | 1 min 31.240 sec | 27.0 | 0.114768 |

| | training_logloss | training_classification_error | validation_rmse \ |
|---|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.124956 | 0.015037 | 0.163222 |
| 2 | 0.088431 | 0.014741 | 0.132450 |
| 3 | 0.075707 | 0.014691 | 0.123747 |

| | | | |
|----|----------|----------|----------|
| 4 | 0.068961 | 0.014758 | 0.120153 |
| 5 | 0.063396 | 0.014701 | 0.118261 |
| 6 | 0.060611 | 0.014688 | 0.117189 |
| 7 | 0.058778 | 0.014663 | 0.116513 |
| 8 | 0.057572 | 0.014613 | 0.116110 |
| 9 | 0.055820 | 0.014637 | 0.115756 |
| 10 | 0.054638 | 0.014632 | 0.115487 |
| 11 | 0.054029 | 0.014639 | 0.115343 |
| 12 | 0.053913 | 0.014633 | 0.115329 |
| 13 | 0.053047 | 0.014637 | 0.115166 |
| 14 | 0.052763 | 0.014629 | 0.115084 |
| 15 | 0.053969 | 0.014625 | 0.115265 |
| 16 | 0.053403 | 0.014623 | 0.115148 |
| 17 | 0.052874 | 0.014616 | 0.115046 |
| 18 | 0.052404 | 0.014608 | 0.114957 |
| 19 | 0.052007 | 0.014606 | 0.114887 |
| 20 | 0.051816 | 0.014607 | 0.114852 |
| 21 | 0.052323 | 0.014605 | 0.114954 |
| 22 | 0.054680 | 0.014605 | 0.115476 |
| 23 | 0.054364 | 0.014602 | 0.115381 |
| 24 | 0.053991 | 0.014596 | 0.115287 |
| 25 | 0.053765 | 0.014590 | 0.115235 |
| 26 | 0.053827 | 0.014593 | 0.115303 |
| 27 | 0.053535 | 0.014593 | 0.115225 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.125809 | 0.015210 |
| 2 | 0.083278 | 0.015008 |
| 3 | 0.070831 | 0.014991 |
| 4 | 0.064559 | 0.014980 |
| 5 | 0.060747 | 0.014986 |
| 6 | 0.058372 | 0.014963 |
| 7 | 0.056646 | 0.014980 |
| 8 | 0.056152 | 0.014963 |
| 9 | 0.055038 | 0.014986 |
| 10 | 0.054155 | 0.014986 |
| 11 | 0.053803 | 0.014991 |
| 12 | 0.053995 | 0.014991 |
| 13 | 0.053385 | 0.014991 |
| 14 | 0.053170 | 0.014986 |
| 15 | 0.054281 | 0.014986 |
| 16 | 0.053809 | 0.014986 |
| 17 | 0.053352 | 0.014991 |
| 18 | 0.052935 | 0.014986 |
| 19 | 0.052648 | 0.014986 |
| 20 | 0.052435 | 0.014986 |
| 21 | 0.052937 | 0.014980 |

| | | |
|----|----------|----------|
| 22 | 0.055460 | 0.014980 |
| 23 | 0.055132 | 0.014986 |
| 24 | 0.054756 | 0.014986 |
| 25 | 0.054538 | 0.015002 |
| 26 | 0.054601 | 0.014991 |
| 27 | 0.054313 | 0.014991 |

Parse progress: || 100%

Parse progress: || 100%

drf Model Build progress: || 100%

Model Details

=====

H2ORandomForestEstimator : Distributed Random Forest

Model Key: rf_K

Model Summary:

| | number_of_trees | number_of_internal_trees | model_size_in_bytes | \ |
|---|-----------------|--------------------------|---------------------|---|
| 0 | 18.0 | 72.0 | 253844.0 | |

| | min_depth | max_depth | mean_depth | min_leaves | max_leaves | mean_leaves |
|---|-----------|-----------|------------|------------|------------|-------------|
| 0 | 8.0 | 17.0 | 12.222222 | 19.0 | 301.0 | 135.45833 |

ModelMetricsMultinomial: drf

** Reported on train data. **

MSE: 0.007878494355481895

RMSE: 0.08876088302558675

LogLoss: 0.03252895137395644

Mean Per-Class Error: 0.14416619061137875

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|----------|--------|----------|---------|----------|-------------------|
| 0 | 301102.0 | 0.0 | 4935.0 | 0.0 | 0.016126 | 4,935 / 306,037 |
| 1 | 0.0 | 2613.0 | 2960.0 | 324.0 | 0.556893 | 3,284 / 5,897 |
| 2 | 926.0 | 16.0 | 671086.0 | 0.0 | 0.001402 | 942 / 672,028 |
| 3 | 0.0 | 2.0 | 180.0 | 80916.0 | 0.002244 | 182 / 81,098 |
| 4 | 302028.0 | 2631.0 | 679161.0 | 81240.0 | 0.008772 | 9,343 / 1,065,060 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.991228 |
| 1 | 2 | 0.999744 |
| 2 | 3 | 0.999992 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: drf
 ** Reported on validation data. **

MSE: 0.0075037751187233
 RMSE: 0.08662433329453854
 LogLoss: 0.03159136187403362
 Mean Per-Class Error: 0.14318723460345475

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|---------|-------|----------|---------|----------|-----------------|
| 0 | 74772.0 | 0.0 | 1189.0 | 0.0 | 0.015653 | 1,189 / 75,961 |
| 1 | 0.0 | 629.0 | 710.0 | 71.0 | 0.553901 | 781 / 1,410 |
| 2 | 266.0 | 2.0 | 168119.0 | 0.0 | 0.001592 | 268 / 168,387 |
| 3 | 0.0 | 1.0 | 32.0 | 20542.0 | 0.001604 | 33 / 20,575 |
| 4 | 75038.0 | 632.0 | 170050.0 | 20613.0 | 0.008527 | 2,271 / 266,333 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.991473 |
| 1 | 2 | 0.999748 |
| 2 | 3 | 0.999989 |
| 3 | 4 | 1.000000 |

Scoring History:

| | timestamp | duration | number_of_trees | training_rmse | \ |
|---|---------------------|------------|-----------------|---------------|-----|
| 0 | 2020-04-04 23:07:10 | 0.018 sec | 0.0 | | NaN |
| 1 | 2020-04-04 23:07:12 | 2.050 sec | 1.0 | 0.103695 | |
| 2 | 2020-04-04 23:07:14 | 3.844 sec | 2.0 | 0.095218 | |
| 3 | 2020-04-04 23:07:16 | 6.100 sec | 3.0 | 0.092595 | |
| 4 | 2020-04-04 23:07:18 | 8.430 sec | 4.0 | 0.091148 | |
| 5 | 2020-04-04 23:07:21 | 10.810 sec | 5.0 | 0.090288 | |

| | | | | |
|----|---------------------|------------|------|----------|
| 6 | 2020-04-04 23:07:24 | 13.550 sec | 6.0 | 0.089748 |
| 7 | 2020-04-04 23:07:26 | 16.278 sec | 7.0 | 0.089364 |
| 8 | 2020-04-04 23:07:29 | 19.189 sec | 8.0 | 0.089033 |
| 9 | 2020-04-04 23:07:32 | 22.120 sec | 9.0 | 0.088683 |
| 10 | 2020-04-04 23:07:35 | 25.293 sec | 10.0 | 0.088459 |
| 11 | 2020-04-04 23:07:39 | 28.614 sec | 11.0 | 0.088302 |
| 12 | 2020-04-04 23:07:42 | 32.331 sec | 12.0 | 0.089790 |
| 13 | 2020-04-04 23:07:46 | 35.909 sec | 13.0 | 0.089455 |
| 14 | 2020-04-04 23:07:50 | 39.794 sec | 14.0 | 0.089164 |
| 15 | 2020-04-04 23:07:54 | 44.109 sec | 15.0 | 0.089371 |
| 16 | 2020-04-04 23:07:58 | 48.403 sec | 16.0 | 0.089119 |
| 17 | 2020-04-04 23:08:03 | 52.771 sec | 17.0 | 0.088915 |
| 18 | 2020-04-04 23:08:07 | 57.139 sec | 18.0 | 0.088761 |

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.051007 | 0.008943 | 0.102997 |
| 2 | 0.040597 | 0.008827 | 0.092417 |
| 3 | 0.036993 | 0.008882 | 0.089348 |
| 4 | 0.035597 | 0.008855 | 0.088099 |
| 5 | 0.034112 | 0.008850 | 0.087375 |
| 6 | 0.033708 | 0.008828 | 0.087004 |
| 7 | 0.033000 | 0.008837 | 0.086788 |
| 8 | 0.032466 | 0.008828 | 0.086622 |
| 9 | 0.031747 | 0.008813 | 0.086492 |
| 10 | 0.031275 | 0.008811 | 0.086418 |
| 11 | 0.031022 | 0.008824 | 0.086340 |
| 12 | 0.033337 | 0.008801 | 0.087080 |
| 13 | 0.032987 | 0.008783 | 0.086933 |
| 14 | 0.032583 | 0.008779 | 0.086789 |
| 15 | 0.033356 | 0.008776 | 0.086876 |
| 16 | 0.033082 | 0.008769 | 0.086763 |
| 17 | 0.032760 | 0.008769 | 0.086687 |
| 18 | 0.032529 | 0.008772 | 0.086624 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.049310 | 0.008794 |
| 2 | 0.035531 | 0.008613 |
| 3 | 0.032043 | 0.008598 |
| 4 | 0.031082 | 0.008591 |
| 5 | 0.030484 | 0.008523 |
| 6 | 0.030347 | 0.008512 |
| 7 | 0.030338 | 0.008523 |
| 8 | 0.030145 | 0.008516 |
| 9 | 0.029882 | 0.008523 |
| 10 | 0.029681 | 0.008527 |
| 11 | 0.029493 | 0.008523 |

| | | |
|----|----------|----------|
| 12 | 0.032044 | 0.008531 |
| 13 | 0.031729 | 0.008527 |
| 14 | 0.031474 | 0.008527 |
| 15 | 0.032308 | 0.008519 |
| 16 | 0.032066 | 0.008516 |
| 17 | 0.031816 | 0.008519 |
| 18 | 0.031591 | 0.008527 |

Variable Importances:

| | variable | relative_importance | scaled_importance | percentage |
|---|-----------|---------------------|-------------------|------------|
| 0 | PD_DESC | 7.348236e+06 | 1.000000 | 0.997538 |
| 1 | PERP_SEX | 8.802860e+03 | 0.001198 | 0.001195 |
| 2 | AGE_GROUP | 5.206346e+03 | 0.000709 | 0.000707 |
| 3 | PERP_RACE | 4.124638e+03 | 0.000561 | 0.000560 |

| | timestamp | duration | number_of_trees | training_rmse | \ |
|----|---------------------|------------|-----------------|---------------|---|
| 0 | 2020-04-04 23:07:10 | 0.018 sec | 0.0 | NaN | |
| 1 | 2020-04-04 23:07:12 | 2.050 sec | 1.0 | 0.103695 | |
| 2 | 2020-04-04 23:07:14 | 3.844 sec | 2.0 | 0.095218 | |
| 3 | 2020-04-04 23:07:16 | 6.100 sec | 3.0 | 0.092595 | |
| 4 | 2020-04-04 23:07:18 | 8.430 sec | 4.0 | 0.091148 | |
| 5 | 2020-04-04 23:07:21 | 10.810 sec | 5.0 | 0.090288 | |
| 6 | 2020-04-04 23:07:24 | 13.550 sec | 6.0 | 0.089748 | |
| 7 | 2020-04-04 23:07:26 | 16.278 sec | 7.0 | 0.089364 | |
| 8 | 2020-04-04 23:07:29 | 19.189 sec | 8.0 | 0.089033 | |
| 9 | 2020-04-04 23:07:32 | 22.120 sec | 9.0 | 0.088683 | |
| 10 | 2020-04-04 23:07:35 | 25.293 sec | 10.0 | 0.088459 | |
| 11 | 2020-04-04 23:07:39 | 28.614 sec | 11.0 | 0.088302 | |
| 12 | 2020-04-04 23:07:42 | 32.331 sec | 12.0 | 0.089790 | |
| 13 | 2020-04-04 23:07:46 | 35.909 sec | 13.0 | 0.089455 | |
| 14 | 2020-04-04 23:07:50 | 39.794 sec | 14.0 | 0.089164 | |
| 15 | 2020-04-04 23:07:54 | 44.109 sec | 15.0 | 0.089371 | |
| 16 | 2020-04-04 23:07:58 | 48.403 sec | 16.0 | 0.089119 | |
| 17 | 2020-04-04 23:08:03 | 52.771 sec | 17.0 | 0.088915 | |
| 18 | 2020-04-04 23:08:07 | 57.139 sec | 18.0 | 0.088761 | |

| | training_logloss | training_classification_error | validation_rmse | \ |
|---|------------------|-------------------------------|-----------------|---|
| 0 | NaN | NaN | NaN | |
| 1 | 0.051007 | 0.008943 | 0.102997 | |
| 2 | 0.040597 | 0.008827 | 0.092417 | |
| 3 | 0.036993 | 0.008882 | 0.089348 | |
| 4 | 0.035597 | 0.008855 | 0.088099 | |

| | | | |
|----|----------|----------|----------|
| 5 | 0.034112 | 0.008850 | 0.087375 |
| 6 | 0.033708 | 0.008828 | 0.087004 |
| 7 | 0.033000 | 0.008837 | 0.086788 |
| 8 | 0.032466 | 0.008828 | 0.086622 |
| 9 | 0.031747 | 0.008813 | 0.086492 |
| 10 | 0.031275 | 0.008811 | 0.086418 |
| 11 | 0.031022 | 0.008824 | 0.086340 |
| 12 | 0.033337 | 0.008801 | 0.087080 |
| 13 | 0.032987 | 0.008783 | 0.086933 |
| 14 | 0.032583 | 0.008779 | 0.086789 |
| 15 | 0.033356 | 0.008776 | 0.086876 |
| 16 | 0.033082 | 0.008769 | 0.086763 |
| 17 | 0.032760 | 0.008769 | 0.086687 |
| 18 | 0.032529 | 0.008772 | 0.086624 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.049310 | 0.008794 |
| 2 | 0.035531 | 0.008613 |
| 3 | 0.032043 | 0.008598 |
| 4 | 0.031082 | 0.008591 |
| 5 | 0.030484 | 0.008523 |
| 6 | 0.030347 | 0.008512 |
| 7 | 0.030338 | 0.008523 |
| 8 | 0.030145 | 0.008516 |
| 9 | 0.029882 | 0.008523 |
| 10 | 0.029681 | 0.008527 |
| 11 | 0.029493 | 0.008523 |
| 12 | 0.032044 | 0.008531 |
| 13 | 0.031729 | 0.008527 |
| 14 | 0.031474 | 0.008527 |
| 15 | 0.032308 | 0.008519 |
| 16 | 0.032066 | 0.008516 |
| 17 | 0.031816 | 0.008519 |
| 18 | 0.031591 | 0.008527 |

Parse progress: || 100%

Parse progress: || 100%

drf Model Build progress: || 100%

Model Details

=====

H2ORandomForestEstimator : Distributed Random Forest

Model Key: rf_M

Model Summary:

| | | | |
|-----------------|--------------------------|---------------------|---|
| number_of_trees | number_of_internal_trees | model_size_in_bytes | \ |
|-----------------|--------------------------|---------------------|---|

| | | | |
|---|------|------|----------|
| 0 | 21.0 | 84.0 | 313167.0 |
|---|------|------|----------|

| | min_depth | max_depth | mean_depth | min_leaves | max_leaves | mean_leaves |
|---|-----------|-----------|------------|------------|------------|-------------|
| 0 | 8.0 | 15.0 | 12.059524 | 23.0 | 274.0 | 133.40475 |

ModelMetricsMultinomial: drf
 ** Reported on train data. **

MSE: 0.005186149758704385
 RMSE: 0.07201492733249396
 LogLoss: 0.027993786119493488
 Mean Per-Class Error: 0.09385858053836384

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|----------|--------|----------|---------|----------|-------------------|
| 0 | 234508.0 | 1.0 | 3484.0 | 0.0 | 0.014643 | 3,485 / 237,993 |
| 1 | 0.0 | 2124.0 | 881.0 | 313.0 | 0.359855 | 1,194 / 3,318 |
| 2 | 313.0 | 110.0 | 707462.0 | 0.0 | 0.000598 | 423 / 707,885 |
| 3 | 0.0 | 0.0 | 29.0 | 85733.0 | 0.000338 | 29 / 85,762 |
| 4 | 234821.0 | 2235.0 | 711856.0 | 86046.0 | 0.004958 | 5,131 / 1,034,958 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.995042 |
| 1 | 2 | 0.999844 |
| 2 | 3 | 0.999993 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: drf
 ** Reported on validation data. **

MSE: 0.004743203790480386
 RMSE: 0.06887092122572766
 LogLoss: 0.028013911209936484
 Mean Per-Class Error: 0.09571930927064987

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|---------|-------|----------|---------|----------|-----------------|
| 0 | 58688.0 | 0.0 | 901.0 | 0.0 | 0.015120 | 901 / 59,589 |
| 1 | 0.0 | 478.0 | 214.0 | 63.0 | 0.366887 | 277 / 755 |
| 2 | 67.0 | 29.0 | 176930.0 | 0.0 | 0.000542 | 96 / 177,026 |
| 3 | 0.0 | 0.0 | 7.0 | 21381.0 | 0.000327 | 7 / 21,388 |
| 4 | 58755.0 | 507.0 | 178052.0 | 21444.0 | 0.004951 | 1,281 / 258,758 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.995049 |
| 1 | 2 | 0.999729 |
| 2 | 3 | 0.999969 |
| 3 | 4 | 1.000000 |

Scoring History:

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|-----------------|-----------------|-----------------|
| 0 | 2020-04-04 23:08:58 | 0.022 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:09:00 | 2.374 sec | 1.0 | 0.166039 |
| 2 | 2020-04-04 23:09:02 | 4.413 sec | 2.0 | 0.123979 |
| 3 | 2020-04-04 23:09:05 | 6.668 sec | 3.0 | 0.106279 |
| 4 | 2020-04-04 23:09:07 | 8.808 sec | 4.0 | 0.095911 |
| 5 | 2020-04-04 23:09:09 | 11.122 sec | 5.0 | 0.089497 |
| 6 | 2020-04-04 23:09:12 | 13.929 sec | 6.0 | 0.085209 |
| 7 | 2020-04-04 23:09:15 | 16.881 sec | 7.0 | 0.082022 |
| 8 | 2020-04-04 23:09:19 | 20.597 sec | 8.0 | 0.079456 |
| 9 | 2020-04-04 23:09:22 | 23.562 sec | 9.0 | 0.077304 |
| 10 | 2020-04-04 23:09:25 | 27.262 sec | 10.0 | 0.075789 |
| 11 | 2020-04-04 23:09:29 | 31.196 sec | 11.0 | 0.074580 |
| 12 | 2020-04-04 23:09:33 | 34.788 sec | 12.0 | 0.073692 |
| 13 | 2020-04-04 23:09:37 | 38.699 sec | 13.0 | 0.072847 |
| 14 | 2020-04-04 23:09:41 | 43.149 sec | 14.0 | 0.072103 |
| 15 | 2020-04-04 23:09:46 | 47.820 sec | 15.0 | 0.075292 |
| 16 | 2020-04-04 23:09:50 | 52.231 sec | 16.0 | 0.074403 |
| 17 | 2020-04-04 23:09:55 | 57.234 sec | 17.0 | 0.073639 |
| 18 | 2020-04-04 23:10:00 | 1 min 1.649 sec | 18.0 | 0.072963 |
| 19 | 2020-04-04 23:10:04 | 1 min 6.498 sec | 19.0 | 0.072374 |

| | training_logloss | training_classification_error | validation_rmse \ |
|---|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.108808 | 0.005319 | 0.166620 |
| 2 | 0.065191 | 0.005108 | 0.115016 |

| | | | |
|----|----------|----------|----------|
| 3 | 0.050088 | 0.005063 | 0.095141 |
| 4 | 0.043323 | 0.005048 | 0.085312 |
| 5 | 0.038516 | 0.005017 | 0.079697 |
| 6 | 0.035525 | 0.005021 | 0.076265 |
| 7 | 0.033542 | 0.005008 | 0.074058 |
| 8 | 0.031699 | 0.004989 | 0.072515 |
| 9 | 0.029933 | 0.004977 | 0.071375 |
| 10 | 0.028702 | 0.004977 | 0.070537 |
| 11 | 0.027752 | 0.004976 | 0.069908 |
| 12 | 0.027185 | 0.004975 | 0.069485 |
| 13 | 0.026499 | 0.004973 | 0.069099 |
| 14 | 0.025806 | 0.004969 | 0.068774 |
| 15 | 0.030603 | 0.004967 | 0.070317 |
| 16 | 0.029843 | 0.004963 | 0.069909 |
| 17 | 0.029192 | 0.004961 | 0.069556 |
| 18 | 0.028481 | 0.004962 | 0.069253 |
| 19 | 0.027942 | 0.004957 | 0.069000 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.113532 | 0.005318 |
| 2 | 0.065646 | 0.005005 |
| 3 | 0.050401 | 0.004970 |
| 4 | 0.042527 | 0.004966 |
| 5 | 0.037757 | 0.004954 |
| 6 | 0.034504 | 0.004947 |
| 7 | 0.032223 | 0.004947 |
| 8 | 0.030603 | 0.004943 |
| 9 | 0.029191 | 0.004947 |
| 10 | 0.028107 | 0.004947 |
| 11 | 0.027241 | 0.004947 |
| 12 | 0.026671 | 0.004947 |
| 13 | 0.026059 | 0.004947 |
| 14 | 0.025484 | 0.004947 |
| 15 | 0.030435 | 0.004947 |
| 16 | 0.029722 | 0.004947 |
| 17 | 0.029041 | 0.004954 |
| 18 | 0.028420 | 0.004954 |
| 19 | 0.027890 | 0.004947 |

See the whole table with `table.as_data_frame()`

Variable Importances:

| variable | relative_importance | scaled_importance | percentage |
|----------|---------------------|-------------------|------------|
|----------|---------------------|-------------------|------------|

| | | | | |
|---|-----------|--------------|----------|----------|
| 0 | PD_DESC | 7.800080e+06 | 1.000000 | 0.997014 |
| 1 | PERP_SEX | 1.110029e+04 | 0.001423 | 0.001419 |
| 2 | PERP_RACE | 8.960558e+03 | 0.001149 | 0.001145 |
| 3 | AGE_GROUP | 3.299111e+03 | 0.000423 | 0.000422 |

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|------------------|-----------------|-----------------|
| 0 | 2020-04-04 23:08:58 | 0.022 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:09:00 | 2.374 sec | 1.0 | 0.166039 |
| 2 | 2020-04-04 23:09:02 | 4.413 sec | 2.0 | 0.123979 |
| 3 | 2020-04-04 23:09:05 | 6.668 sec | 3.0 | 0.106279 |
| 4 | 2020-04-04 23:09:07 | 8.808 sec | 4.0 | 0.095911 |
| 5 | 2020-04-04 23:09:09 | 11.122 sec | 5.0 | 0.089497 |
| 6 | 2020-04-04 23:09:12 | 13.929 sec | 6.0 | 0.085209 |
| 7 | 2020-04-04 23:09:15 | 16.881 sec | 7.0 | 0.082022 |
| 8 | 2020-04-04 23:09:19 | 20.597 sec | 8.0 | 0.079456 |
| 9 | 2020-04-04 23:09:22 | 23.562 sec | 9.0 | 0.077304 |
| 10 | 2020-04-04 23:09:25 | 27.262 sec | 10.0 | 0.075789 |
| 11 | 2020-04-04 23:09:29 | 31.196 sec | 11.0 | 0.074580 |
| 12 | 2020-04-04 23:09:33 | 34.788 sec | 12.0 | 0.073692 |
| 13 | 2020-04-04 23:09:37 | 38.699 sec | 13.0 | 0.072847 |
| 14 | 2020-04-04 23:09:41 | 43.149 sec | 14.0 | 0.072103 |
| 15 | 2020-04-04 23:09:46 | 47.820 sec | 15.0 | 0.075292 |
| 16 | 2020-04-04 23:09:50 | 52.231 sec | 16.0 | 0.074403 |
| 17 | 2020-04-04 23:09:55 | 57.234 sec | 17.0 | 0.073639 |
| 18 | 2020-04-04 23:10:00 | 1 min 1.649 sec | 18.0 | 0.072963 |
| 19 | 2020-04-04 23:10:04 | 1 min 6.498 sec | 19.0 | 0.072374 |
| 20 | 2020-04-04 23:10:09 | 1 min 11.361 sec | 20.0 | 0.071883 |
| 21 | 2020-04-04 23:10:15 | 1 min 17.112 sec | 21.0 | 0.072015 |

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.108808 | 0.005319 | 0.166620 |
| 2 | 0.065191 | 0.005108 | 0.115016 |
| 3 | 0.050088 | 0.005063 | 0.095141 |
| 4 | 0.043323 | 0.005048 | 0.085312 |
| 5 | 0.038516 | 0.005017 | 0.079697 |
| 6 | 0.035525 | 0.005021 | 0.076265 |
| 7 | 0.033542 | 0.005008 | 0.074058 |
| 8 | 0.031699 | 0.004989 | 0.072515 |
| 9 | 0.029933 | 0.004977 | 0.071375 |
| 10 | 0.028702 | 0.004977 | 0.070537 |
| 11 | 0.027752 | 0.004976 | 0.069908 |
| 12 | 0.027185 | 0.004975 | 0.069485 |
| 13 | 0.026499 | 0.004973 | 0.069099 |
| 14 | 0.025806 | 0.004969 | 0.068774 |
| 15 | 0.030603 | 0.004967 | 0.070317 |

| | | | |
|----|----------|----------|----------|
| 16 | 0.029843 | 0.004963 | 0.069909 |
| 17 | 0.029192 | 0.004961 | 0.069556 |
| 18 | 0.028481 | 0.004962 | 0.069253 |
| 19 | 0.027942 | 0.004957 | 0.069000 |
| 20 | 0.027438 | 0.004959 | 0.068784 |
| 21 | 0.027994 | 0.004958 | 0.068871 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.113532 | 0.005318 |
| 2 | 0.065646 | 0.005005 |
| 3 | 0.050401 | 0.004970 |
| 4 | 0.042527 | 0.004966 |
| 5 | 0.037757 | 0.004954 |
| 6 | 0.034504 | 0.004947 |
| 7 | 0.032223 | 0.004947 |
| 8 | 0.030603 | 0.004943 |
| 9 | 0.029191 | 0.004947 |
| 10 | 0.028107 | 0.004947 |
| 11 | 0.027241 | 0.004947 |
| 12 | 0.026671 | 0.004947 |
| 13 | 0.026059 | 0.004947 |
| 14 | 0.025484 | 0.004947 |
| 15 | 0.030435 | 0.004947 |
| 16 | 0.029722 | 0.004947 |
| 17 | 0.029041 | 0.004954 |
| 18 | 0.028420 | 0.004954 |
| 19 | 0.027890 | 0.004947 |
| 20 | 0.027406 | 0.004947 |
| 21 | 0.028014 | 0.004951 |

Parse progress: || 100%

Parse progress: || 100%

drf Model Build progress: || 100%

Model Details

=====

H2ORandomForestEstimator : Distributed Random Forest

Model Key: rf_S

Model Summary:

| | number_of_trees | number_of_internal_trees | model_size_in_bytes | \ |
|---|-----------------|--------------------------|---------------------|---|
| 0 | 72.0 | 288.0 | 559503.0 | |

| | min_depth | max_depth | mean_depth | min_leaves | max_leaves | mean_leaves |
|---|-----------|-----------|------------|------------|------------|-------------|
| 0 | 4.0 | 16.0 | 11.017361 | 7.0 | 176.0 | 77.21875 |

ModelMetricsMultinomial: drf
** Reported on train data. **

MSE: 0.007346465050413919
RMSE: 0.08571152227334385
LogLoss: 0.0384756856978329
Mean Per-Class Error: 0.10515733899416524

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|---------|-------|---------|-------|----------|-----------------|
| 0 | 38371.0 | 0.0 | 766.0 | 0.0 | 0.019572 | 766 / 39,137 |
| 1 | 3.0 | 260.0 | 154.0 | 8.0 | 0.388235 | 165 / 425 |
| 2 | 112.0 | 8.0 | 92113.0 | 0.0 | 0.001301 | 120 / 92,233 |
| 3 | 4.0 | 0.0 | 6.0 | 858.0 | 0.011521 | 10 / 868 |
| 4 | 38490.0 | 268.0 | 93039.0 | 866.0 | 0.007998 | 1,061 / 132,663 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.992002 |
| 1 | 2 | 0.999812 |
| 2 | 3 | 0.999955 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: drf
** Reported on validation data. **

MSE: 0.007450859906088194
RMSE: 0.08631836366665087
LogLoss: 0.03921656968285433
Mean Per-Class Error: 0.14390466445553032

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|--------|------|---------|-------|----------|-------------|
| 0 | 9516.0 | 0.0 | 191.0 | 0.0 | 0.019677 | 191 / 9,707 |
| 1 | 0.0 | 45.0 | 51.0 | 4.0 | 0.550000 | 55 / 100 |
| 2 | 16.0 | 1.0 | 23150.0 | 0.0 | 0.000734 | 17 / 23,167 |
| 3 | 1.0 | 0.0 | 0.0 | 191.0 | 0.005208 | 1 / 192 |

4 9533.0 46.0 23392.0 195.0 0.007960 264 / 33,166

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.992040 |
| 1 | 2 | 0.999879 |
| 2 | 3 | 0.999940 |
| 3 | 4 | 1.000000 |

Scoring History:

| | timestamp | duration | number_of_trees | training_rmse | \ |
|----|---------------------|------------|-----------------|---------------|---|
| 0 | 2020-04-04 23:10:55 | 0.056 sec | 0.0 | NaN | |
| 1 | 2020-04-04 23:10:55 | 0.296 sec | 1.0 | 0.226961 | |
| 2 | 2020-04-04 23:10:56 | 0.561 sec | 2.0 | 0.165825 | |
| 3 | 2020-04-04 23:10:56 | 0.849 sec | 3.0 | 0.140145 | |
| 4 | 2020-04-04 23:10:56 | 1.282 sec | 4.0 | 0.125124 | |
| 5 | 2020-04-04 23:10:57 | 1.796 sec | 5.0 | 0.114856 | |
| 6 | 2020-04-04 23:10:58 | 2.497 sec | 6.0 | 0.108226 | |
| 7 | 2020-04-04 23:10:58 | 3.192 sec | 7.0 | 0.104147 | |
| 8 | 2020-04-04 23:10:59 | 3.829 sec | 8.0 | 0.100035 | |
| 9 | 2020-04-04 23:11:00 | 4.449 sec | 9.0 | 0.096980 | |
| 10 | 2020-04-04 23:11:00 | 5.151 sec | 10.0 | 0.094922 | |
| 11 | 2020-04-04 23:11:01 | 5.863 sec | 11.0 | 0.093135 | |
| 12 | 2020-04-04 23:11:02 | 6.513 sec | 12.0 | 0.096629 | |
| 13 | 2020-04-04 23:11:02 | 7.263 sec | 13.0 | 0.094786 | |
| 14 | 2020-04-04 23:11:03 | 7.982 sec | 14.0 | 0.093312 | |
| 15 | 2020-04-04 23:11:04 | 8.749 sec | 15.0 | 0.092905 | |
| 16 | 2020-04-04 23:11:05 | 9.535 sec | 16.0 | 0.091924 | |
| 17 | 2020-04-04 23:11:06 | 10.380 sec | 17.0 | 0.090973 | |
| 18 | 2020-04-04 23:11:06 | 11.271 sec | 18.0 | 0.090195 | |
| 19 | 2020-04-04 23:11:07 | 12.177 sec | 19.0 | 0.089566 | |

| | training_logloss | training_classification_error | validation_rmse | \ |
|---|------------------|-------------------------------|-----------------|---|
| 0 | NaN | NaN | NaN | |
| 1 | 0.219991 | 0.047796 | 0.227545 | |
| 2 | 0.130328 | 0.023896 | 0.142502 | |
| 3 | 0.102013 | 0.016387 | 0.116691 | |
| 4 | 0.089853 | 0.012812 | 0.105161 | |
| 5 | 0.074680 | 0.010814 | 0.098635 | |
| 6 | 0.065567 | 0.009814 | 0.095051 | |
| 7 | 0.062841 | 0.009299 | 0.092862 | |

| | | | |
|----|----------|----------|----------|
| 8 | 0.056350 | 0.008776 | 0.091060 |
| 9 | 0.051707 | 0.008500 | 0.089832 |
| 10 | 0.049436 | 0.008367 | 0.088979 |
| 11 | 0.047377 | 0.008238 | 0.088421 |
| 12 | 0.052354 | 0.008182 | 0.090589 |
| 13 | 0.050047 | 0.008080 | 0.089795 |
| 14 | 0.048098 | 0.008027 | 0.089139 |
| 15 | 0.047328 | 0.008000 | 0.089227 |
| 16 | 0.046395 | 0.008004 | 0.088834 |
| 17 | 0.044568 | 0.007993 | 0.088408 |
| 18 | 0.043259 | 0.007977 | 0.088025 |
| 19 | 0.042452 | 0.007984 | 0.087722 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.223925 | 0.049056 |
| 2 | 0.110307 | 0.008382 |
| 3 | 0.084242 | 0.008231 |
| 4 | 0.069287 | 0.008201 |
| 5 | 0.058296 | 0.008171 |
| 6 | 0.054095 | 0.008171 |
| 7 | 0.050323 | 0.008171 |
| 8 | 0.047905 | 0.008141 |
| 9 | 0.044297 | 0.008141 |
| 10 | 0.042820 | 0.008141 |
| 11 | 0.041729 | 0.008141 |
| 12 | 0.047622 | 0.008141 |
| 13 | 0.046254 | 0.008141 |
| 14 | 0.045070 | 0.008141 |
| 15 | 0.045226 | 0.008081 |
| 16 | 0.044318 | 0.008081 |
| 17 | 0.043466 | 0.008081 |
| 18 | 0.042697 | 0.008081 |
| 19 | 0.042069 | 0.008020 |

See the whole table with `table.as_data_frame()`

Variable Importances:

| | variable | relative_importance | scaled_importance | percentage |
|---|-----------|---------------------|-------------------|------------|
| 0 | PD_DESC | 3.048465e+06 | 1.000000 | 0.996221 |
| 1 | PERP_RACE | 7.735905e+03 | 0.002538 | 0.002528 |
| 2 | PERP_SEX | 2.010547e+03 | 0.000660 | 0.000657 |
| 3 | AGE_GROUP | 1.817694e+03 | 0.000596 | 0.000594 |

| | timestamp | duration | number_of_trees | training_rmse | \ |
|-------|---------------------|------------------|-----------------|---------------|---|
| 0 | 2020-04-04 23:10:55 | 0.056 sec | 0.0 | NaN | |
| 1 | 2020-04-04 23:10:55 | 0.296 sec | 1.0 | 0.226961 | |
| 2 | 2020-04-04 23:10:56 | 0.561 sec | 2.0 | 0.165825 | |
| 3 | 2020-04-04 23:10:56 | 0.849 sec | 3.0 | 0.140145 | |
| 4 | 2020-04-04 23:10:56 | 1.282 sec | 4.0 | 0.125124 | |
| 5 | 2020-04-04 23:10:57 | 1.796 sec | 5.0 | 0.114856 | |
| 6 | 2020-04-04 23:10:58 | 2.497 sec | 6.0 | 0.108226 | |
| 7 | 2020-04-04 23:10:58 | 3.192 sec | 7.0 | 0.104147 | |
| 8 | 2020-04-04 23:10:59 | 3.829 sec | 8.0 | 0.100035 | |
| 9 | 2020-04-04 23:11:00 | 4.449 sec | 9.0 | 0.096980 | |
| 10 | 2020-04-04 23:11:00 | 5.151 sec | 10.0 | 0.094922 | |
| 11 | 2020-04-04 23:11:01 | 5.863 sec | 11.0 | 0.093135 | |
| 12 | 2020-04-04 23:11:02 | 6.513 sec | 12.0 | 0.096629 | |
| 13 | 2020-04-04 23:11:02 | 7.263 sec | 13.0 | 0.094786 | |
| 14 | 2020-04-04 23:11:03 | 7.982 sec | 14.0 | 0.093312 | |
| 15 | 2020-04-04 23:11:04 | 8.749 sec | 15.0 | 0.092905 | |
| 16 | 2020-04-04 23:11:05 | 9.535 sec | 16.0 | 0.091924 | |
| 17 | 2020-04-04 23:11:06 | 10.380 sec | 17.0 | 0.090973 | |
| 18 | 2020-04-04 23:11:06 | 11.271 sec | 18.0 | 0.090195 | |
| 19 | 2020-04-04 23:11:07 | 12.177 sec | 19.0 | 0.089566 | |
| 20 | 2020-04-04 23:11:08 | 13.154 sec | 20.0 | 0.088963 | |
| 21 | 2020-04-04 23:11:09 | 14.204 sec | 21.0 | 0.088840 | |
| 22 | 2020-04-04 23:11:11 | 15.338 sec | 22.0 | 0.089836 | |
| 23 | 2020-04-04 23:11:12 | 16.520 sec | 23.0 | 0.089381 | |
| 24 | 2020-04-04 23:11:13 | 17.777 sec | 24.0 | 0.088905 | |
| 25 | 2020-04-04 23:11:14 | 19.139 sec | 25.0 | 0.088586 | |
| 26 | 2020-04-04 23:11:16 | 20.431 sec | 26.0 | 0.088806 | |
| 27 | 2020-04-04 23:11:17 | 21.793 sec | 27.0 | 0.088492 | |
| 28 | 2020-04-04 23:11:18 | 23.319 sec | 28.0 | 0.088252 | |
| 29 | 2020-04-04 23:11:20 | 24.847 sec | 29.0 | 0.087899 | |
| | ... | ... | ... | ... | |
| 43 | 2020-04-04 23:11:49 | 54.095 sec | 43.0 | 0.086432 | |
| 44 | 2020-04-04 23:11:52 | 56.453 sec | 44.0 | 0.086342 | |
| 45 | 2020-04-04 23:11:54 | 58.809 sec | 45.0 | 0.086196 | |
| 46 | 2020-04-04 23:11:56 | 1 min 1.230 sec | 46.0 | 0.086126 | |
| 47 | 2020-04-04 23:11:59 | 1 min 4.074 sec | 47.0 | 0.085994 | |
| 48 | 2020-04-04 23:12:02 | 1 min 6.671 sec | 48.0 | 0.085878 | |
| 49 | 2020-04-04 23:12:04 | 1 min 9.210 sec | 49.0 | 0.085822 | |
| 50 | 2020-04-04 23:12:07 | 1 min 11.674 sec | 50.0 | 0.085776 | |
| 51 | 2020-04-04 23:12:09 | 1 min 14.219 sec | 51.0 | 0.085716 | |
| 52 | 2020-04-04 23:12:12 | 1 min 16.817 sec | 52.0 | 0.085619 | |
| 53 | 2020-04-04 23:12:15 | 1 min 19.419 sec | 53.0 | 0.085674 | |
| 54 | 2020-04-04 23:12:17 | 1 min 22.098 sec | 54.0 | 0.085575 | |
| 55 | 2020-04-04 23:12:21 | 1 min 25.810 sec | 55.0 | 0.085489 | |
| 56 | 2020-04-04 23:12:25 | 1 min 29.489 sec | 56.0 | 0.085446 | |
| 57 | 2020-04-04 23:12:28 | 1 min 33.021 sec | 57.0 | 0.085402 | |

| | | | | |
|----|---------------------|------------------|------|----------|
| 58 | 2020-04-04 23:12:31 | 1 min 35.983 sec | 58.0 | 0.085330 |
| 59 | 2020-04-04 23:12:34 | 1 min 38.876 sec | 59.0 | 0.085260 |
| 60 | 2020-04-04 23:12:37 | 1 min 41.838 sec | 60.0 | 0.085199 |
| 61 | 2020-04-04 23:12:40 | 1 min 44.953 sec | 61.0 | 0.085130 |
| 62 | 2020-04-04 23:12:43 | 1 min 48.074 sec | 62.0 | 0.085073 |
| 63 | 2020-04-04 23:12:47 | 1 min 51.375 sec | 63.0 | 0.085061 |
| 64 | 2020-04-04 23:12:50 | 1 min 54.607 sec | 64.0 | 0.085033 |
| 65 | 2020-04-04 23:12:53 | 1 min 57.988 sec | 65.0 | 0.085126 |
| 66 | 2020-04-04 23:12:57 | 2 min 1.391 sec | 66.0 | 0.085480 |
| 67 | 2020-04-04 23:13:00 | 2 min 4.909 sec | 67.0 | 0.085414 |
| 68 | 2020-04-04 23:13:04 | 2 min 8.418 sec | 68.0 | 0.085732 |
| 69 | 2020-04-04 23:13:08 | 2 min 12.362 sec | 69.0 | 0.085720 |
| 70 | 2020-04-04 23:13:11 | 2 min 16.243 sec | 70.0 | 0.085655 |
| 71 | 2020-04-04 23:13:16 | 2 min 20.442 sec | 71.0 | 0.085585 |
| 72 | 2020-04-04 23:13:19 | 2 min 24.075 sec | 72.0 | 0.085712 |

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.219991 | 0.047796 | 0.227545 |
| 2 | 0.130328 | 0.023896 | 0.142502 |
| 3 | 0.102013 | 0.016387 | 0.116691 |
| 4 | 0.089853 | 0.012812 | 0.105161 |
| 5 | 0.074680 | 0.010814 | 0.098635 |
| 6 | 0.065567 | 0.009814 | 0.095051 |
| 7 | 0.062841 | 0.009299 | 0.092862 |
| 8 | 0.056350 | 0.008776 | 0.091060 |
| 9 | 0.051707 | 0.008500 | 0.089832 |
| 10 | 0.049436 | 0.008367 | 0.088979 |
| 11 | 0.047377 | 0.008238 | 0.088421 |
| 12 | 0.052354 | 0.008182 | 0.090589 |
| 13 | 0.050047 | 0.008080 | 0.089795 |
| 14 | 0.048098 | 0.008027 | 0.089139 |
| 15 | 0.047328 | 0.008000 | 0.089227 |
| 16 | 0.046395 | 0.008004 | 0.088834 |
| 17 | 0.044568 | 0.007993 | 0.088408 |
| 18 | 0.043259 | 0.007977 | 0.088025 |
| 19 | 0.042452 | 0.007984 | 0.087722 |
| 20 | 0.041777 | 0.007976 | 0.087448 |
| 21 | 0.041698 | 0.007945 | 0.087403 |
| 22 | 0.044085 | 0.007998 | 0.088103 |
| 23 | 0.043571 | 0.008013 | 0.087859 |
| 24 | 0.042986 | 0.007975 | 0.087593 |
| 25 | 0.042521 | 0.007983 | 0.087426 |
| 26 | 0.043367 | 0.008013 | 0.087668 |
| 27 | 0.042514 | 0.008043 | 0.087548 |
| 28 | 0.042086 | 0.008043 | 0.087419 |
| 29 | 0.041386 | 0.008035 | 0.087227 |
| .. | ... | ... | ... |

| | | | |
|----|----------|----------|----------|
| 43 | 0.038706 | 0.008035 | 0.086537 |
| 44 | 0.038488 | 0.008043 | 0.086484 |
| 45 | 0.038275 | 0.008058 | 0.086420 |
| 46 | 0.038073 | 0.008066 | 0.086393 |
| 47 | 0.037857 | 0.008050 | 0.086324 |
| 48 | 0.037644 | 0.008020 | 0.086266 |
| 49 | 0.037469 | 0.008035 | 0.086235 |
| 50 | 0.037311 | 0.008050 | 0.086210 |
| 51 | 0.037144 | 0.008073 | 0.086187 |
| 52 | 0.036964 | 0.008050 | 0.086125 |
| 53 | 0.037193 | 0.008050 | 0.086175 |
| 54 | 0.037008 | 0.008050 | 0.086104 |
| 55 | 0.036840 | 0.008035 | 0.086050 |
| 56 | 0.036712 | 0.008035 | 0.086029 |
| 57 | 0.036548 | 0.008035 | 0.086003 |
| 58 | 0.036415 | 0.008035 | 0.085975 |
| 59 | 0.036272 | 0.008028 | 0.085932 |
| 60 | 0.036151 | 0.007998 | 0.085883 |
| 61 | 0.036007 | 0.007983 | 0.085839 |
| 62 | 0.035878 | 0.007975 | 0.085802 |
| 63 | 0.035813 | 0.007990 | 0.085798 |
| 64 | 0.035699 | 0.007983 | 0.085776 |
| 65 | 0.036153 | 0.007983 | 0.085839 |
| 66 | 0.037270 | 0.007990 | 0.086103 |
| 67 | 0.037127 | 0.007990 | 0.086064 |
| 68 | 0.038020 | 0.007990 | 0.086310 |
| 69 | 0.038018 | 0.008005 | 0.086304 |
| 70 | 0.037869 | 0.007990 | 0.086261 |
| 71 | 0.037726 | 0.007983 | 0.086220 |
| 72 | 0.038476 | 0.007998 | 0.086318 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.223925 | 0.049056 |
| 2 | 0.110307 | 0.008382 |
| 3 | 0.084242 | 0.008231 |
| 4 | 0.069287 | 0.008201 |
| 5 | 0.058296 | 0.008171 |
| 6 | 0.054095 | 0.008171 |
| 7 | 0.050323 | 0.008171 |
| 8 | 0.047905 | 0.008141 |
| 9 | 0.044297 | 0.008141 |
| 10 | 0.042820 | 0.008141 |
| 11 | 0.041729 | 0.008141 |
| 12 | 0.047622 | 0.008141 |
| 13 | 0.046254 | 0.008141 |
| 14 | 0.045070 | 0.008141 |
| 15 | 0.045226 | 0.008081 |

| | | |
|----|----------|----------|
| 16 | 0.044318 | 0.008081 |
| 17 | 0.043466 | 0.008081 |
| 18 | 0.042697 | 0.008081 |
| 19 | 0.042069 | 0.008020 |
| 20 | 0.041465 | 0.007930 |
| 21 | 0.041768 | 0.007960 |
| 22 | 0.043419 | 0.008081 |
| 23 | 0.042888 | 0.007990 |
| 24 | 0.042291 | 0.007960 |
| 25 | 0.041864 | 0.007960 |
| 26 | 0.042682 | 0.007960 |
| 27 | 0.042289 | 0.008020 |
| 28 | 0.041837 | 0.008111 |
| 29 | 0.041383 | 0.007960 |
| .. | ... | ... |
| 43 | 0.039306 | 0.008020 |
| 44 | 0.039081 | 0.008020 |
| 45 | 0.038902 | 0.008020 |
| 46 | 0.038708 | 0.008020 |
| 47 | 0.038495 | 0.008020 |
| 48 | 0.038300 | 0.007960 |
| 49 | 0.038137 | 0.007990 |
| 50 | 0.037976 | 0.008020 |
| 51 | 0.037829 | 0.008020 |
| 52 | 0.037646 | 0.008020 |
| 53 | 0.037897 | 0.008020 |
| 54 | 0.037721 | 0.008020 |
| 55 | 0.037560 | 0.007960 |
| 56 | 0.037426 | 0.008020 |
| 57 | 0.037280 | 0.008020 |
| 58 | 0.037139 | 0.007990 |
| 59 | 0.037001 | 0.007960 |
| 60 | 0.036860 | 0.007960 |
| 61 | 0.036723 | 0.007960 |
| 62 | 0.036596 | 0.007960 |
| 63 | 0.036532 | 0.007960 |
| 64 | 0.036409 | 0.007960 |
| 65 | 0.036860 | 0.007960 |
| 66 | 0.037969 | 0.007960 |
| 67 | 0.037834 | 0.007960 |
| 68 | 0.038707 | 0.007960 |
| 69 | 0.038698 | 0.007960 |
| 70 | 0.038558 | 0.007960 |
| 71 | 0.038423 | 0.007960 |
| 72 | 0.039217 | 0.007960 |

[73 rows x 10 columns]
Parse progress: || 100%

Parse progress: || 100%
 drf Model Build progress: || 100%
 Model Details
 =====
 H2ORandomForestEstimator : Distributed Random Forest
 Model Key: rf_B

Model Summary:

| | number_of_trees | number_of_internal_trees | model_size_in_bytes | \ |
|---|-----------------|--------------------------|---------------------|---|
| 0 | 18.0 | 72.0 | 225684.0 | |

| | min_depth | max_depth | mean_depth | min_leaves | max_leaves | mean_leaves |
|---|-----------|-----------|------------|------------|------------|-------------|
| 0 | 6.0 | 16.0 | 11.444445 | 18.0 | 281.0 | 114.388885 |

ModelMetricsMultinomial: drf
 ** Reported on train data. **

MSE: 0.004301819027950368
 RMSE: 0.06558825373457025
 LogLoss: 0.01762195615523751
 Mean Per-Class Error: 0.0763673557544016

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|----------|-------|----------|---------|----------|-----------------|
| 0 | 218978.0 | 0.0 | 3555.0 | 0.0 | 0.015975 | 3,555 / 222,533 |
| 1 | 0.0 | 855.0 | 283.0 | 61.0 | 0.286906 | 344 / 1,199 |
| 2 | 280.0 | 36.0 | 632586.0 | 0.0 | 0.000499 | 316 / 632,902 |
| 3 | 0.0 | 1.0 | 33.0 | 16240.0 | 0.002089 | 34 / 16,274 |
| 4 | 219258.0 | 892.0 | 636457.0 | 16301.0 | 0.004868 | 4,249 / 872,908 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.995132 |
| 1 | 2 | 0.999968 |
| 2 | 3 | 0.999992 |
| 3 | 4 | 1.000000 |

ModelMetricsMultinomial: drf
 ** Reported on validation data. **

MSE: 0.004327694745189184
 RMSE: 0.0657852167678209
 LogLoss: 0.017743214393310787
 Mean Per-Class Error: 0.08434717159805417

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|---------|-------|----------|--------|----------|-----------------|
| 0 | 54689.0 | 0.0 | 909.0 | 0.0 | 0.016350 | 909 / 55,598 |
| 1 | 0.0 | 207.0 | 84.0 | 13.0 | 0.319079 | 97 / 304 |
| 2 | 72.0 | 10.0 | 158137.0 | 0.0 | 0.000518 | 82 / 158,219 |
| 3 | 0.0 | 1.0 | 5.0 | 4155.0 | 0.001442 | 6 / 4,161 |
| 4 | 54761.0 | 218.0 | 159135.0 | 4168.0 | 0.005012 | 1,094 / 218,282 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.994988 |
| 1 | 2 | 0.999973 |
| 2 | 3 | 0.999995 |
| 3 | 4 | 1.000000 |

Scoring History:

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|------------|-----------------|-----------------|
| 0 | 2020-04-04 23:14:10 | 0.032 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:14:12 | 1.663 sec | 1.0 | 0.067871 |
| 2 | 2020-04-04 23:14:14 | 3.524 sec | 2.0 | 0.066201 |
| 3 | 2020-04-04 23:14:16 | 5.655 sec | 3.0 | 0.065632 |
| 4 | 2020-04-04 23:14:18 | 7.584 sec | 4.0 | 0.065524 |
| 5 | 2020-04-04 23:14:20 | 9.667 sec | 5.0 | 0.065255 |
| 6 | 2020-04-04 23:14:22 | 11.958 sec | 6.0 | 0.065056 |
| 7 | 2020-04-04 23:14:24 | 14.322 sec | 7.0 | 0.065896 |
| 8 | 2020-04-04 23:14:27 | 17.144 sec | 8.0 | 0.065649 |
| 9 | 2020-04-04 23:14:30 | 19.809 sec | 9.0 | 0.065457 |
| 10 | 2020-04-04 23:14:33 | 22.789 sec | 10.0 | 0.065246 |
| 11 | 2020-04-04 23:14:36 | 25.878 sec | 11.0 | 0.065151 |
| 12 | 2020-04-04 23:14:39 | 29.337 sec | 12.0 | 0.065361 |

| | | | | |
|----|---------------------|------------|------|----------|
| 13 | 2020-04-04 23:14:43 | 32.769 sec | 13.0 | 0.065231 |
| 14 | 2020-04-04 23:14:46 | 36.291 sec | 14.0 | 0.065119 |
| 15 | 2020-04-04 23:14:50 | 40.176 sec | 15.0 | 0.065969 |
| 16 | 2020-04-04 23:14:54 | 44.274 sec | 16.0 | 0.065860 |
| 17 | 2020-04-04 23:14:58 | 48.253 sec | 17.0 | 0.065706 |
| 18 | 2020-04-04 23:15:02 | 52.359 sec | 18.0 | 0.065588 |

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.028939 | 0.005201 | 0.068914 |
| 2 | 0.022993 | 0.005039 | 0.066048 |
| 3 | 0.020637 | 0.004975 | 0.065604 |
| 4 | 0.019811 | 0.004963 | 0.065573 |
| 5 | 0.018931 | 0.004939 | 0.065458 |
| 6 | 0.018174 | 0.004911 | 0.065472 |
| 7 | 0.018541 | 0.004902 | 0.065921 |
| 8 | 0.018006 | 0.004895 | 0.065786 |
| 9 | 0.017551 | 0.004886 | 0.065689 |
| 10 | 0.017039 | 0.004874 | 0.065639 |
| 11 | 0.016750 | 0.004871 | 0.065582 |
| 12 | 0.017699 | 0.004873 | 0.065683 |
| 13 | 0.017518 | 0.004866 | 0.065628 |
| 14 | 0.017173 | 0.004861 | 0.065588 |
| 15 | 0.018110 | 0.004865 | 0.065948 |
| 16 | 0.017946 | 0.004872 | 0.065898 |
| 17 | 0.017757 | 0.004868 | 0.065840 |
| 18 | 0.017622 | 0.004868 | 0.065785 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.031050 | 0.005388 |
| 2 | 0.019038 | 0.005117 |
| 3 | 0.016551 | 0.005044 |
| 4 | 0.016561 | 0.005053 |
| 5 | 0.016279 | 0.005021 |
| 6 | 0.016074 | 0.005021 |
| 7 | 0.016771 | 0.005016 |
| 8 | 0.016735 | 0.005012 |
| 9 | 0.016599 | 0.004998 |
| 10 | 0.016431 | 0.005007 |
| 11 | 0.016350 | 0.005012 |
| 12 | 0.017345 | 0.005012 |
| 13 | 0.017236 | 0.005012 |
| 14 | 0.017111 | 0.005012 |
| 15 | 0.018171 | 0.005012 |
| 16 | 0.018028 | 0.005012 |
| 17 | 0.017888 | 0.005012 |
| 18 | 0.017743 | 0.005012 |

Variable Importances:

| | variable | relative_importance | scaled_importance | percentage |
|---|-----------|---------------------|-------------------|------------|
| 0 | PD_DESC | 4.853892e+06 | 1.000000 | 0.998627 |
| 1 | PERP_SEX | 2.881350e+03 | 0.000594 | 0.000593 |
| 2 | AGE_GROUP | 2.620734e+03 | 0.000540 | 0.000539 |
| 3 | PERP_RACE | 1.169469e+03 | 0.000241 | 0.000241 |

| | timestamp | duration | number_of_trees | training_rmse \ |
|----|---------------------|------------|-----------------|-----------------|
| 0 | 2020-04-04 23:14:10 | 0.032 sec | 0.0 | NaN |
| 1 | 2020-04-04 23:14:12 | 1.663 sec | 1.0 | 0.067871 |
| 2 | 2020-04-04 23:14:14 | 3.524 sec | 2.0 | 0.066201 |
| 3 | 2020-04-04 23:14:16 | 5.655 sec | 3.0 | 0.065632 |
| 4 | 2020-04-04 23:14:18 | 7.584 sec | 4.0 | 0.065524 |
| 5 | 2020-04-04 23:14:20 | 9.667 sec | 5.0 | 0.065255 |
| 6 | 2020-04-04 23:14:22 | 11.958 sec | 6.0 | 0.065056 |
| 7 | 2020-04-04 23:14:24 | 14.322 sec | 7.0 | 0.065896 |
| 8 | 2020-04-04 23:14:27 | 17.144 sec | 8.0 | 0.065649 |
| 9 | 2020-04-04 23:14:30 | 19.809 sec | 9.0 | 0.065457 |
| 10 | 2020-04-04 23:14:33 | 22.789 sec | 10.0 | 0.065246 |
| 11 | 2020-04-04 23:14:36 | 25.878 sec | 11.0 | 0.065151 |
| 12 | 2020-04-04 23:14:39 | 29.337 sec | 12.0 | 0.065361 |
| 13 | 2020-04-04 23:14:43 | 32.769 sec | 13.0 | 0.065231 |
| 14 | 2020-04-04 23:14:46 | 36.291 sec | 14.0 | 0.065119 |
| 15 | 2020-04-04 23:14:50 | 40.176 sec | 15.0 | 0.065969 |
| 16 | 2020-04-04 23:14:54 | 44.274 sec | 16.0 | 0.065860 |
| 17 | 2020-04-04 23:14:58 | 48.253 sec | 17.0 | 0.065706 |
| 18 | 2020-04-04 23:15:02 | 52.359 sec | 18.0 | 0.065588 |

| | training_logloss | training_classification_error | validation_rmse \ |
|----|------------------|-------------------------------|-------------------|
| 0 | NaN | NaN | NaN |
| 1 | 0.028939 | 0.005201 | 0.068914 |
| 2 | 0.022993 | 0.005039 | 0.066048 |
| 3 | 0.020637 | 0.004975 | 0.065604 |
| 4 | 0.019811 | 0.004963 | 0.065573 |
| 5 | 0.018931 | 0.004939 | 0.065458 |
| 6 | 0.018174 | 0.004911 | 0.065472 |
| 7 | 0.018541 | 0.004902 | 0.065921 |
| 8 | 0.018006 | 0.004895 | 0.065786 |
| 9 | 0.017551 | 0.004886 | 0.065689 |
| 10 | 0.017039 | 0.004874 | 0.065639 |
| 11 | 0.016750 | 0.004871 | 0.065582 |

| | | | |
|----|----------|----------|----------|
| 12 | 0.017699 | 0.004873 | 0.065683 |
| 13 | 0.017518 | 0.004866 | 0.065628 |
| 14 | 0.017173 | 0.004861 | 0.065588 |
| 15 | 0.018110 | 0.004865 | 0.065948 |
| 16 | 0.017946 | 0.004872 | 0.065898 |
| 17 | 0.017757 | 0.004868 | 0.065840 |
| 18 | 0.017622 | 0.004868 | 0.065785 |

| | validation_logloss | validation_classification_error |
|----|--------------------|---------------------------------|
| 0 | NaN | NaN |
| 1 | 0.031050 | 0.005388 |
| 2 | 0.019038 | 0.005117 |
| 3 | 0.016551 | 0.005044 |
| 4 | 0.016561 | 0.005053 |
| 5 | 0.016279 | 0.005021 |
| 6 | 0.016074 | 0.005021 |
| 7 | 0.016771 | 0.005016 |
| 8 | 0.016735 | 0.005012 |
| 9 | 0.016599 | 0.004998 |
| 10 | 0.016431 | 0.005007 |
| 11 | 0.016350 | 0.005012 |
| 12 | 0.017345 | 0.005012 |
| 13 | 0.017236 | 0.005012 |
| 14 | 0.017111 | 0.005012 |
| 15 | 0.018171 | 0.005012 |
| 16 | 0.018028 | 0.005012 |
| 17 | 0.017888 | 0.005012 |
| 18 | 0.017743 | 0.005012 |

Here, we notice a relatively accurate model with an overall error of a mere 1.4593%, but a significant error in the prediction of Infractions.

Note that when we included the law code as a predictor (not shown in report), this error of infractions was also eliminated in a random forest model, result in a negligibly small error on all classifications. The following was the output obtained:

Parse progress: || 100% Parse progress: || 100% drf Model Build progress: || 100%

4 Model Details

H2ORandomForestEstimator : Distributed Random Forest Model Key: rf_covType_v1

| Model Summary: | number_of_trees | number_of_internal_trees | model_size_in_bytes |
|--|-----------------|--------------------------|---------------------|
| min_depth max_depth mean_depth min_leaves max_leaves mean_leaves | -- | ----- | --- |
| | | | 51 204 |

533899 4 18 11.5882 7 237 80.7157

ModelMetricsMultinomial: drf ** Reported on train data. **

MSE: 0.0001734852476826297 RMSE: 0.013171379870105854 LogLoss: 0.003255583616165607
Mean Per-Class Error: 0.00033000525330283565

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class F I M V Error Rate ----- 513090 0 28 0 5.45683e-05 28 / 513,118 0 7954 8 2 0.00125565 10 / 7,964 1 1 1.21166e+06 0 1.65062e-06 2 / 1,211,665 1 0 0 122674 8.15162e-06 1 / 122,675 513092 7955 1.2117e+06 122676 2.20974e-05 41 / 1,855,422

Top-4 Hit Ratios: k hit_ratio --- 1 0.999978 2 0.999999 3 1 4 1

ModelMetricsMultinomial: drf ** Reported on validation data. **

MSE: 0.00016996249120756354 RMSE: 0.013036966334525972 LogLoss: 0.0035461151687810766 Mean Per-Class Error: 0.0017149228257738725

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class F I M V Error Rate ----- 127951 0 14 0 0.000109405 14 / 127,965 0 2061 14 0 0.00674699 14 / 2,075 0 1 303173 0 3.29844e-06 1 / 303,174 0 0 0 30642 0 0 / 30,642 127951 2062 303201 30642 6.25194e-05 29 / 463,856

Top-4 Hit Ratios: k hit_ratio --- 1 0.999938 2 0.999963 3 0.999998 4 1

Variable Importances: variable relative_importance scaled_importance percentage ----- PD_DESC 3.08297e+07 1 0.853075 LAW_CODE 5.21138e+06 0.169038 0.144202 ARREST_BORO 50738.3 0.00164576 0.00140396 AGE_GROUP 27099.1 0.000878995 0.000749848 PERP_SEX 20594.9 0.00066802 0.000569871

Finally, we test the model trained on the Bronx borough (notorious for being relatively unsafe) against the data for the Manhattan borough (the most prestigious borough). If there were a bias, we would expect to see a lot of misclassifications, skewed towards more serious levels of offenses like Felony or Violation as opposed to Infraction or Misdemeanor.

```
In [53]: saved_model = h2o.load_model("C:/Users/varch/Desktop/Github/NYPD-Data-Comp/rf_B/rf_B")
```

```
In [54]: testingset = h2o.H2OFrame(nypd.loc[nypd['ARREST_BORO'] == 'M'])
        saved_model.model_performance(testingset)
```

Parse progress: || 100%

ModelMetricsMultinomial: drf
** Reported on test data. **

MSE: 0.0048710613504178765
RMSE: 0.06979298926409354
LogLoss: 0.023246667485524727
Mean Per-Class Error: 0.10366541412714929

Confusion Matrix: Row labels: Actual class; Column labels: Predicted class

| | F | I | M | V | Error | Rate |
|---|----------|--------|----------|----------|----------|-------------------|
| 0 | 293200.0 | 0.0 | 4404.0 | 0.0 | 0.014798 | 4,404 / 297,604 |
| 1 | 0.0 | 2458.0 | 1243.0 | 372.0 | 0.396514 | 1,615 / 4,073 |
| 2 | 387.0 | 75.0 | 884499.0 | 0.0 | 0.000522 | 462 / 884,961 |
| 3 | 0.0 | 180.0 | 123.0 | 106848.0 | 0.002828 | 303 / 107,151 |
| 4 | 293587.0 | 2713.0 | 890269.0 | 107220.0 | 0.005244 | 6,784 / 1,293,789 |

Top-4 Hit Ratios:

| | k | hit_ratio |
|---|---|-----------|
| 0 | 1 | 0.994756 |
| 1 | 2 | 0.999790 |
| 2 | 3 | 0.999930 |
| 3 | 4 | 1.000000 |

Out [54] :

However, we still get great predictions with an overall error rate of just 0.5244%

5 Conclusions

From all of our modeling and data exploration, we can conclude that the data does not represent discrimination on the basis of ethnographic variables. We reconcile this notion with the understanding that the data set is simply unable to capture the instances where discrimination occurs. Our recommendation for an accurate analysis of discrimination is to additionally capture two sets of occurrences - before and after the arrests:

1. The stopping point before an arrest is actually made on the suspicion of illegal activity (arbitrary profiling)
2. The varying harshness of judgements given to individuals having committed the same crime but belonging to different demographics

Data representing these factors, when incorporated into our existing random forest model, is very likely to discern any patterns of discrimination present, considering the extreme accuracy with which our model currently provides predictions based on the provided data.